

HyperLogLog: Analysis and implementation of an improved algorithm

Dequeker Chloé, Ziat Ghiles

February 2015

Contents

1	Introduction	3
1.1	cardinality estimation problem	3
2	LinearCounting and HyperLogLog	3
2.1	LinearCounting	3
2.2	HyperLogLog	3
3	HyperLogLog++	4
3.1	transition to 64 bits	4
3.2	Bias estimation and correction	4
3.3	Memory optimization	5
3.3.1	Sparse representation	5
3.3.2	Dense representation	5
3.3.3	Varint encoding	5
4	Conclusion	5

1 Introduction

In this paper, we present our implementation and analysis of a cardinality estimation algorithm proposed by Stefan Heule, Marc Nunkesser and Alexander Hall: **HyperLogLog++**. This algorithm is itself an improvement of the **HyperLogLog** algorithm proposed by Flajolet et. al.

1.1 cardinality estimation problem

Finding the number of distinct elements in a data set with duplicates is a well-known problem which applies in many fields.

The naive solution to this problem is to examine for each element of the data stream its belonging to a data structure \mathcal{D} . If \mathcal{D} does not contain the element we add it to the data structure. At the end of the process, the cardinality of the data stream is equal to the size of \mathcal{D} .

This solution gives the exact answer but it is easy to see that it scales very badly as the size of the data stream grows.

In order to resolve this problem, several algorithms have been proposed. These include **LinearCounting** and **HyperLogLog** which are the two bases of the studied algorithm.

2 LinearCounting and HyperLogLog

2.1 LinearCounting

2.2 HyperLogLog

The approach of the **HyperLogLog** algorithm to approximate the cardinalities of a multiset is completely different. It is based on randomization using a hash function for each element of the multiset. It then focuses on the maximum of the number of leading zeros in each hash values. It is legitimate to expect that the more items there will be, the more this value will be high. To improve the precision of this calculation, **HyperLogLog** uses the stochastic averaging technique: Doing so, it splits the stream in m substreams, and perform the computation separately on each.

The result is then subjected to corrections:

- *Small range correction* : As shown by simulation, for a cardinality smaller than $\frac{5}{2}$ of the number of substreams, non-linear distortions appear. For that range, **LinearCounting** is used.
- *Large range correction* : Due to the use of a 32 bit hash function, when the cardinality goes to 2^{32} , the chances of hash collisions increases.

3 HyperLogLog++

3.1 transition to 64 bits

Using a 32-bits hash function restricts the area of efficiency of the algorithm to the sets with less then 2^{32} distincts elements. That's why an proposed improvement is to use a 64-bits hash function. It does not significantly change the memory cost (it is only increased only by 1 bit per substream).

3.2 Bias estimation and correction

For a given configuration of the algorithm, the observed bias is only dependant on the cardinality estimated. From this observation, we implement a correction method: As shown in figure 1, the raw estimation of HLL is distorted for small cardinalities. In order to correct this error, we take measures of it for cardinalities between 0 and 100 000 (with a step of 500) and we store them into a file. From now, the file will be loaded at the begining of the calculation. A correction may then be calculated for the result using a linear interpolation between the values registered and the raw estimations.




Figure 1: Cardinality estimation for the corrected HLL

3.3 Memory optimization

Memory usage is an important factor in order to have a efficient algorithm. In this algorithm, we can see that the size of the different values we are using don't need to be of the standard size of an int (which would be 4 bytes for a 32 bits integer). We need in fact to keep two size of values, the first one being the index which maximum value is 2^P with P the precision factor. That means any index could be stocked on 14 bits. The second value is the number of leading 0 of the hashed value which can't be over $64 - P$ bits since we work on a 64 bit version. The result is that the number of leading 0 will need 6 bits at most. The total size of those two values is then of $6 + P = 20$ bits. We will show in the next sections the different kinds of compression we used during the implementation and in the final state of the algorithm.

3.3.1 Sparse representation

This is the first type of compression, and is the one which should be used when only a low number of index have been hashed. This representation works by pairs (index, number of leading zero (clz)). For a better understanding consider a bitmap, then separate it by 20 bits blocks. Each of these block will be a pair (index, value). The first P bits of the pair will represent the value of the index, and the next 6 bits the value of clz. The total size (for $P = 14$) of the bitmap will then be the number of different index hashed times 20 bits. We can then easily see why this representation is particularly efficient for a low number of indexes and this is it's strong perk. On the other hand, the more the number of different indexes grow, the less efficient this representation becomes.

We will then introduce the Dense representation, which becomes more efficient when the number of index reaches a certain value we will be talking about in the next section.

3.3.2 Dense representation

We will introduce in this section the second type of compression. As we said earlier this representation is more efficient with a high number of indexes. In this representation, we need here to divide the bitmap by 6bits blocks. In this representation, when we go through the first 6 bits of the bitmap, we will read the value of index 0. The next 6 bits after that will be the value of index 1 and so on. This representation allows us to represent the pair (index,value) without writing the index. We can easily see this bitmap will be of constant size since the value of the index is deducted from the position of the 6-bit block in the bitmap.

Considering those two representation, it is clear we need to start the algorithm using the sparse representation, and then switch at some point to the dense one. The limit where we want to switch from one to another is when the sparse representation take more memory than the dense one. In the implementation, it is then important to keep track of the bitmap's size and switch to the dense representation whenever it is necessary.

3.3.3 Varint encoding

Since the temporary set used in the sparse representation is merged with the list before it gets too large, performing a compression on it is not as interesting then on the sorted list. We'll try to reduce the memory usage of it by playing on two points:

- Using fixed-size integers as it is common practice in many languages may here result in a waste of memory space.
- Since the manipulated list is sorted, we can take advantage of this information.

4 Conclusion